

Interpreting logistic models

Frank Edwards

2/27/2023

- Remember that logistic regression is a GLM with a logit link function

- Remember that logistic regression is a GLM with a logit link function
- A GLM takes the form: $g(y) = \mathbf{X}\beta$

- Remember that logistic regression is a GLM with a logit link function
- A GLM takes the form: $g(y) = \mathbf{X}\beta$
- Logistic regression is the special case where g is the logit function:
 $\text{logit}(y) = \mathbf{X}\beta$

- Remember that logistic regression is a GLM with a logit link function
- A GLM takes the form: $g(y) = \mathbf{X}\beta$
- Logistic regression is the special case where g is the logit function:
 $\text{logit}(y) = \mathbf{X}\beta$
- A logistic regression model returns $\mathbf{X}\beta$ on the logit scale

- Remember that logistic regression is a GLM with a logit link function
- A GLM takes the form: $g(y) = \mathbf{X}\beta$
- Logistic regression is the special case where g is the logit function:
 $\text{logit}(y) = \mathbf{X}\beta$
- A logistic regression model returns $\mathbf{X}\beta$ on the logit scale
- How can we convert $\mathbf{x}\beta$ to something useful?

Let's return to the grad school admission example

```
admits <- read_csv("./data/binary.csv")  
summary(admits)
```

##	admit	gre	gpa	rank
##	Min. :0.0000	Min. :220.0	Min. :2.260	Min. :1.000
##	1st Qu.:0.0000	1st Qu.:520.0	1st Qu.:3.130	1st Qu.:2.000
##	Median :0.0000	Median :580.0	Median :3.395	Median :2.000
##	Mean :0.3175	Mean :587.7	Mean :3.390	Mean :2.485
##	3rd Qu.:1.0000	3rd Qu.:660.0	3rd Qu.:3.670	3rd Qu.:3.000
##	Max. :1.0000	Max. :800.0	Max. :4.000	Max. :4.000

Let's look at this as the distribution of the probability of admissions across the data

- First, fit an intercept-only logistic regression model

```
m0 <- glm(admit ~ 1, data = admits, family = "binomial")  
m0_est <- tidy(m0)
```

- What does this model tell us?

What does this model tell us?

```
m0_est$estimate ## log odds
```

```
## [1] -0.7652847
```

```
exp(m0_est$estimate) ## odds
```

```
## [1] 0.4652015
```

```
exp(m0_est$estimate)/(1 + exp(m0_est$estimate)) ## probability
```

```
## [1] 0.3175
```

```
mean(admits$admit) ## mean admission probability
```

```
## [1] 0.3175
```



Let's add a predictor

```
m1 <- glm(admit ~ 1 + gre, data = admits, family = "binomial")
tidy(m1)
```

```
## # A tibble: 2 x 5
```

##	term	estimate	std.error	statistic	p.value
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	(Intercept)	-2.90	0.606	-4.79	0.00000169
## 2	gre	0.00358	0.000986	3.63	0.000280

Predictions for all data points in M0 and M1 - what's going on?



Two linear predictors

Why do these generate such different predictions?



Intercept only model (m0): $p(admit) = \beta_0$

With GRE predictor (m1): $p(admit) = \beta_0 + \beta_1 GRE_i$

To ease interpretation, let's scale GRE

Scale mean-centers and SD scales variables: $\text{scale}(x_i) = \frac{x_i - \bar{x}}{sd(x)}$

Linear transformations of variables: mean-center and SD scale



Re-estimate the model: much nicer to look at

```
admits <- admits %>%  
  mutate(gre_s = as.numeric(scale(gre)))  
m1 <- glm(admit ~ 1 + gre_s, data = admits, family = "binomial")  
m1_est <- tidy(m1)  
m1_est
```

```
## # A tibble: 2 x 5  
##   term          estimate std.error statistic  p.value  
##   <chr>          <dbl>     <dbl>     <dbl>    <dbl>  
## 1 (Intercept)  -0.796     0.111     -7.20 6.01e-13  
## 2 gre_s        0.414     0.114      3.63 2.80e- 4
```

Interpret the model

```
m1_est
```

```
## # A tibble: 2 x 5
```

```
##   term          estimate std.error statistic  p.value  
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>  
## 1 (Intercept)   -0.796     0.111     -7.20 6.01e-13  
## 2 gre_s         0.414     0.114      3.63 2.80e- 4
```

Remember: $\text{logit}(y) = X\beta = \log\left(\frac{y}{1-y}\right)$

So: $y = \text{logit}^{-1}(X\beta) = \frac{\exp(X\beta)}{\exp(X\beta)+1}$

Interpret the model

```
m1_est
```

```
## # A tibble: 2 x 5
```

```
##   term          estimate std.error statistic  p.value  
##   <chr>         <dbl>     <dbl>     <dbl>    <dbl>  
## 1 (Intercept)  -0.796      0.111      -7.20 6.01e-13  
## 2 gre_s        0.414      0.114       3.63 2.80e- 4
```

Remember: $\text{logit}(y) = X\beta = \log\left(\frac{y}{1-y}\right)$

So: $y = \text{logit}^{-1}(X\beta) = \frac{\exp(X\beta)}{\exp(X\beta)+1}$

- What is β_0 ?

Interpret the model

```
m1_est
```

```
## # A tibble: 2 x 5
```

```
##   term          estimate std.error statistic  p.value  
##   <chr>         <dbl>     <dbl>     <dbl>    <dbl>  
## 1 (Intercept)  -0.796      0.111      -7.20 6.01e-13  
## 2 gre_s         0.414      0.114       3.63 2.80e- 4
```

Remember: $\text{logit}(y) = X\beta = \log\left(\frac{y}{1-y}\right)$

So: $y = \text{logit}^{-1}(X\beta) = \frac{\exp(X\beta)}{\exp(X\beta)+1}$

- What is β_0 ?
- What is β_1 ?

$$e^{y_1+y_2} = e^{y_1}e^{y_2}$$

Refresher on exponentials

$$e^{y_1+y_2} = e^{y_1}e^{y_2}$$

and

$$e^{y_1-y_2} = \frac{e^{y_1}}{e^{y_2}}$$

Refresher on exponentials

$$e^{y_1+y_2} = e^{y_1} e^{y_2}$$

and

$$e^{y_1-y_2} = \frac{e^{y_1}}{e^{y_2}}$$

so how can we rewrite:

$$\exp(\text{logit}(y)) = \frac{y}{1-y} = e^{\beta_0 + \beta_1 x_1}$$

On the log scale, β_0 and β_1 are related to y multiplicatively because

$$e^{\beta_0 + \beta_1 x_1} = e^{\beta_0} e^{\beta_1 x_1}$$

Odds are defined as the probability of the event occurring divided by the probability of probability of the event not occurring. To obtain odds in a logistic regression, we exponentiate both sides:

$$\frac{y}{1-y} = e^{\beta_0 + \beta_1 x_1}$$

Odds are defined as the probability of the event occurring divided by the probability of probability of the event not occurring. To obtain odds in a logistic regression, we exponentiate both sides:

$$\frac{y}{1-y} = e^{\beta_0 + \beta_1 x_1}$$

The odds of $y = 1$ are simply $e^{x\beta}$

The odds ratio is the ratio of two odds - or the proportional change in odds. We can obtain an isolated estimate for the relationship between $\beta_1 x_{1i}$ and y this way:

$$\frac{\text{Odds}(y|x_1 = 1)}{\text{Odds}(y|x_1 = 0)} = \frac{e^{x\beta + \beta_1}}{e^{x\beta}} = \frac{e^{x\beta} \times e^{\beta_1}}{e^{x\beta}} = e^{\beta_1}$$

The odds ratio can be interpreted as the change in odds of $y = 1$ for a one-unit change in x_1 .

- Odds ratios appear convenient - e^{β_1} is a percent change in y for a one-unit change in x_1

- Odds ratios appear convenient - e^{β_1} is a percent change in y for a one-unit change in x_1

How do they work?

In our example: what do these figures mean?

```
new_dat <- c(1, 0) # for scale(gre) == 0, mean score
odds_0 <- exp(new_dat %*% m1_est$estimate)
odds_0
```

```
##           [,1]
## [1,] 0.4510945
```

```
new_dat1 <- c(1, 1)
odds_1 <- exp(new_dat1 %*% m1_est$estimate)
odds_1
```

```
##           [,1]
## [1,] 0.6823082
```

```
odds_1/odds_0 # odds ratio
```

```
##           [,1]
## [1,] 1.512562
```

```
exp(m1_est$estimate[2]) # exp(beta_1)
```

```
## [1] 1.512562
```

The odds of admission are 1.5125617 times higher for a student with a GRE score one standard deviation above the mean than they are for a student with a mean GRE score.

Interpreting the odds ratio

The odds of admission are 1.5125617 times higher for a student with a GRE score one standard deviation above the mean than they are for a student with a mean GRE score.


Any trouble you can anticipate here?

Interpreting the odds ratio

The odds of admission are 1.5125617 times higher for a student with a GRE score one standard deviation above the mean than they are for a student with a mean GRE score.

Any trouble you can anticipate here?

A visual example: the “effect” of 1 SD increase in GRE scores on $\Pr(\text{admit}=1)$




It is easy enough to work on the probability scale

To obtain predicted probabilities of the observed:

- `p_hat<-predict(m1, type = "response")`

On the probability scale

```
preds <- predict(m1, type = "response")  
p_hat <- data.frame(gre = admits$gre, p = preds)  
  
ggplot(p_hat, aes(x = gre, y = p)) + geom_line()
```



The basic logic of prediction

1. Choose scenarios of theoretical interest
2. Define these in terms of “counterfactual” (fake) data
3. Plug these fake data into the linear predictor (regression equation)
4. Visualize!

The basic logic of prediction

Reminder: our model is

$$\text{logit}(p(\text{admit}_i)) = \beta_0 + \beta_1 \text{GRE}_i$$

$$\text{admit}_i \sim \text{Binomial}(1, p)$$

```
m1 <- glm(admit ~ gre, data = admits, family = "binomial")
```

1. Choose scenarios of theoretical interest

Low GRE, average GRE, high GRE

Define these scenarios in R

2. Define these in terms of “counterfactual” (fake) data

```
## Look at the distribution of the data to think about scenarios
```

```
mean(admits$gre)
```

```
## [1] 587.7
```

```
sd(admits$gre)
```

```
## [1] 115.5165
```

Define these scenarios in R

2. Define these in terms of “counterfactual” (fake) data

```
## Look at the distribution of the data to think about scenarios
```

```
mean(admits$gre)
```

```
## [1] 587.7
```

```
sd(admits$gre)
```

```
## [1] 115.5165
```

Let's define scenarios at the mean, 1 SD below the mean, and 1 SD above the mean

```
fake_data <- data.frame(gre = c(mean(admits$gre), mean(admits$gre) - sd(admits$gre),  
  mean(admits$gre) + sd(admits$gre)))
```

```
fake_data
```

```
##      gre
```

```
## 1 587.7000
```

```
## 2 472.1835
```

```
## 3 703.2165
```

Generating expected probabilities

3. Plug these fake data into the linear predictor (regression equation)

Because $\text{logit}(p(\text{admit}_i)) = \beta_0 + \beta_1 x_i$, we can compute the expected probability of admission for a student with mean GRE scores as

```
coef(m1)
```

```
## (Intercept)          gre  
## -2.901344270  0.003582212
```

```
### mean GRE scenario: linear predictor
```

```
exp(-2.9 + 0.0036 * 587.7)/(1 + exp(-2.9 + 0.0036 * 587.7))
```

```
## [1] 0.3133982
```

Generating expected probabilities

3. Plug these fake data into the linear predictor (regression equation)

The `predict()` function makes life very easy here

```
## linear predictor
```

```
predict(m1, newdata = fake_data)
```

```
##           1           2           3
```

```
## -0.7960785 -1.2098832 -0.3822738
```

```
## probability scale (inverse logit)
```

```
predict(m1, newdata = fake_data, type = "response")
```

```
##           1           2           3
```

```
## 0.3108650 0.2297217 0.4055786
```


4. Visualize!

```
### set up our data frame with predictions for plotting
```

```
fake_data <- fake_data %>%
```

```
  mutate(p_hat = predict(m1, newdata = fake_data, type = "response"))
```

```
ggplot(fake_data, aes(x = gre, y = p_hat)) + geom_point()
```

